

COS 226, SPRING 2014

**ALGORITHMS
AND
DATA STRUCTURES**

KEVIN WAYNE



**PRINCETON
UNIVERSITY**

<http://www.princeton.edu/~cos226>

COS 226 course overview

What is COS 226?

- Intermediate-level survey course.
- Programming and problem solving, with applications.
- **Algorithm:** method for solving a problem.
- **Data structure:** method to store information.

topic	data structures and algorithms
data types	stack, queue, bag, union-find, priority queue
sorting	quicksort, mergesort, heapsort, radix sorts
searching	BST, red-black BST, hash table
graphs	BFS, DFS, Prim, Kruskal, Dijkstra
strings	KMP, regular expressions, tries, data compression
advanced	B-tree, k-d tree, suffix array, maxflow

Why study algorithms?

Their impact is broad and far-reaching.

Internet. Web search, packet routing, distributed file sharing, ...

Biology. Human genome project, protein folding, ...

Computers. Circuit layout, file system, compilers, ...

Computer graphics. Movies, video games, virtual reality, ...

Security. Cell phones, e-commerce, voting machines, ...



Multimedia. MP3, JPG, DivX, HDTV, face recognition, ...

Social networks. Recommendations, news feeds, advertisements, ...

Physics. N-body simulation, particle collision simulation, ...

:

Google
YAHOO!
bing



Why study algorithms?

Their impact is broad and far-reaching.

Mysterious algorithm was 4% of trading activity last week


October 11, 2012

A single mysterious computer program that placed orders — and then subsequently canceled them — made up 4 percent of all quote traffic in the U.S. stock market last week, according to the top tracker of [high-frequency trading](#) activity.

The motive of the algorithm is still unclear, [CNBC](#) reports.

The program placed orders in 25-millisecond bursts involving about 500 stocks, according to Nanex, a market data firm. The algorithm never executed a single trade, and it abruptly ended at about 10:30 a.m. ET Friday.

"My guess is that the algo was testing the market, as high-frequency frequently does," says Jon Najarian, co-founder of TradeMonster.com. "As soon as they add bandwidth, the HFT crowd sees how quickly they can top out to create latency." ([Read More: Unclear What Caused Kraft Spike: Nanex Founder.](#))




Generic high frequency trading chart (credit: Nanex)

Why study algorithms?

Old roots, new opportunities.

- Study of algorithms dates at least to Euclid.
- Formalized by Church and Turing in 1930s.
- Some important algorithms were discovered by undergraduates in a course like this!




Why study algorithms?

For intellectual stimulation.

“For me, great algorithms are the poetry of computation. Just like verse, they can be terse, allusive, dense, and even mysterious. But once unlocked, they cast a brilliant new light on some aspect of computing.” — Francis Sullivan

“ An algorithm must be seen to be believed. ” — Donald Knuth



FROM THE
EDITORS

THE JOY OF ALGORITHMS

A black and white portrait of Francis Sullivan, a man with short, dark hair and a warm smile. He is wearing a light-colored t-shirt. The image is positioned at the top left of the page.

THE THEME OF THIS FIRST-OF-THE-CENTURY ISSUE OF COMPUTING IN SCIENCE & ENGINEERING IS ALGORITHMS. IN FACT, WE WERE BOLD ENOUGH—AND PERHAPS FOOLISH ENOUGH—to call the 10 examples we've selected “THE TOP 10 ALGORITHMS OF THE CENTURY.”

Computational algorithms are probably as old as civilization. Sumerian cuneiform, one of the most ancient written records, consists partly of algorithmic descriptions for reckoning in base 60. And I suppose we could claim that the Droid algorithm for “Estimating the startle hand” is embedded in Starcheese. (That’s really the startle hand?)

Like many other advances that technology affords, algorithms have moved in leaps and sometimes bounds. The 70s and 80s, at least in computing, were maturing. The algorithm was becoming more than a simple task or might not have even been.

Perhaps that’s why a number of people have problems like I do.

The file says to "exercise its mother of invention" is "invention creates its own success". Our need for powerful methods to solve complex problems has led us to believe that there must be a single method that can be applied to all problems. As we have seen, however, the search for such a method has led us to the dead end of NP-complete problems and to an attempt to answer many simple questions. There are three possible ways to proceed:

I suspect that at the 21st century, things will be very far as other revolution in understanding the foundations of mathematics and physics will be needed. The first is quantum computing and problems associated with the generation of random numbers. The second is the theory of the physical world. The third is the theory of computation.

Quantum computers are not going to be very useful, but it

Just like verse, they can be terse, allusive, dense, and even
is not going to be dull either! ☺

— 1 —

Why study algorithms?


To become a proficient programmer.

“I will, in fact, claim that the difference between a bad programmer and a good one is whether he considers his code or his data structures more important. Bad programmers worry about the code. Good programmers worry about data structures and their relationships. ”

— Linus Torvalds (creator of Linux)



“Algorithms + Data Structures = Programs. ” — Niklaus Wirth



Why study algorithms?

They may unlock the secrets of life and of the universe.


“ Computer models mirroring real life have become crucial for most advances made in chemistry today.... Today the computer is just as important a tool for chemists as the test tube. ”

— Royal Swedish Academy of Sciences

(Nobel Prize in Chemistry 2013)




Martin Karplus, Michael Levitt, and Arieh Warshel



Why study algorithms?

To solve problems that could not otherwise be addressed.



http://www.youtube.com/watch?v=ua7YIN4eL_w

Why study algorithms?

Everybody else is doing it.

```
% sort -rn PU2013-14.txt
```

774	COS 126	General Computer Science
615	ECO 100	Introduction to Microeconomics
471	ECO 101	Introduction to Macroeconomics
444	ENG 385	Children's Literature
440	MAT 202	Linear Algebra with Applications
414	COS 226	Algorithms and Data Structures
405	MAT 201	Multivariable Calculus
384	CHV 310	Practical Ethics
344	REL 261	Christian Ethics and Modern Society
320	PSY 101	Introduction to Psychology
300	COS 217	Introduction to Programming Systems
...		

Why study algorithms?

For fun and profit.



Apple Computer

facebook®

CISCO SYSTEMS



Google™

IBM

Nintendo®

JANE STREET

Morgan Stanley

NETFLIX

Adobe™

RSA
SECURITY™

D E Shaw & Co

ORACLE®

P
PANDORA®

Akamai

YAHOO!

amazon.com

Microsoft®

P I X A R
ANIMATION STUDIOS

Why study algorithms?

- Their impact is broad and far-reaching.
- Old roots, new opportunities.
- For intellectual stimulation.
- To become a proficient programmer.
- They may unlock the secrets of life and of the universe.
- To solve problems that could not otherwise be addressed.
- Everybody else is doing it.
- For fun and profit.

Why study anything else?



Lectures

Traditional lectures. Introduce new material.

Electronic devices. Permitted, but only to enhance lecture.



no



no



no

What	When	Where	Who	Office Hours
L01	MW 11-12:20	McCosh 10	Kevin Wayne	see web

Lectures

Traditional lectures. Introduce new material.

Flipped lectures.

- Watch videos online **before** lecture.
- Complete pre-lecture activities.
- Attend only one "flipped" lecture per week (interactive, collaborative, experimental).
- Apply via web ASAP: results by 5pm today.



What	When	Where	Who	Office Hours
L01	MW 11-12:20	McCosh 10	Kevin Wayne	see web
L02	W 11-12:20	Frist 307	Andy Guna Josh Hug	see web

Precepts

Discussion, problem-solving, background for assignments.

What	When	Where	Who	Office Hours
P01	Th 11-11:50	CS 102	Andy Guna †	see web
P02	Th 12:30-1:20	Bobst 105	Andy Guna †	see web
P03	Th 1:30-2:20	Bobst 105	Nevin Li	see web
P04	F 10-10:50	Bobst 105	Jennifer Guo	see web
P05	F 11-11:50	Bobst 105	Madhu Jayakumar	see web
P05A	F 11-11:50	Sherrerd 001	Ruth Dannenfelser	see web
P06	F 2:30-3:20	Friend 108	Chris Eubank	see web
P06A	F 2:30-3:20	Friend 111	TBA	see web
P06B	F 2:30-3:20	Friend 109	Josh Hug †	see web
P07	F 3:30-4:20	Friend 108	Josh Hug †	see web



likely to change

† lead preceptor

Coursework and grading

Programming assignments. 45%

- Due on Tuesdays at 11pm via electronic submission.
- Collaboration/lateness policies: see web.

Exercises. 10%


- Due on Sundays at 11pm in Blackboard.
- Collaboration/lateness policies: see web.

Exams. 15% + 30%

- Midterm (in class on Wednesday, March 12).
- Final (to be scheduled by Registrar).


Staff discretion. [adjust borderline cases]

- Report errata.
- Contribute to Piazza discussion forum.
- Attend and participate in precept/lecture.




Resources (textbook)


Required reading. Algorithms 4th edition by R. Sedgewick and K. Wayne, Addison-Wesley Professional, 2011, ISBN 0-321-57351-X.



1st edition (1982)




2nd edition (1988)



3rd edition (1997)

3rd book scanned
by Google books



4th edition (2011)

Available in hardcover and Kindle.

- Online: Amazon (\$60/\$35 to buy), Chegg (\$25 to rent), ...
- Brick-and-mortar: Labyrinth Books (122 Nassau St).
- On reserve: Engineering library.

Resources (web)

Course content.

- Course info.
- Lecture slides.
- Flipped lectures.
- Programming assignments.
- Exercises.
- Exam archive.

 PRINCETON
UNIVERSITY

COMPUTER SCIENCE 226
ALGORITHMS AND DATA STRUCTURES
SPRING 2014

[Course Information](#) | [Lectures](#) | [Flipped](#) | [Precepts](#) | [Assignments](#) | [Exercises](#) | [Exams](#)


COURSE INFORMATION

Description. This course surveys the most important algorithms and data structures in use on computers today. Particular emphasis is given to algorithms for sorting, searching, and string processing. Fundamental algorithms in a number of other areas are covered as well, including geometric and graph algorithms. The course will concentrate on developing implementations, understanding their performance characteristics, and estimating their potential effectiveness in applications.

<http://www.princeton.edu/~cos226>

Booksite.

- Brief summary of content.
- Download code from book.
- APIs and Javadoc.



ALGORITHMS, 4TH EDITION


*essential information that
every serious programmer
needs to know about
algorithms and data structures*

Textbook. The textbook *Algorithms, 4th Edition* by Robert Sedgewick and Kevin Wayne [[Amazon](#) · [Addison-Wesley](#)] surveys the most important algorithms and data structures in use today. The textbook is organized into six chapters:

- *Chapter 1: Fundamentals* introduces a scientific and engineering basis for comparing algorithms and making predictions. It also includes our programming model.
- *Chapter 2: Sorting* considers several classic sorting algorithms, including insertion sort, mergesort, and quicksort. It also includes a binary heap implementation of a priority queue.
- *Chapter 3: Searching* describes several classic symbol table implementations, including binary search trees, red-black trees, and hash tables.

<http://algs4.cs.princeton.edu>

Resources (web)



<http://www.princeton.edu/~cos226>

Resources (web)

226 – Google Search

http://www.google.com/#sclient=psy&hl=en&q=+226&aq=f&i

226 – Google Search

Web Images Videos Maps News Shopping Gmail more ▾ Web History | Search settings | Sign in

Google

226

About 236,000,000 results (0.18 seconds) Advanced search

Instant is on ▾

Everything

Images Videos News More

Any time Latest Past 2 days

All results Sites with images

More search tools

[Area codes 519 and 226 - Wikipedia, the free encyclopedia](#)
The 226 area code was first proposed as a result of an NPA exhaustion study conducted in the 1990s. The issue was raised with the CRTC by telecommunications ...
en.wikipedia.org/wiki/Area_codes_519_and_226 - Cached - Similar

[226 - Wikipedia, the free encyclopedia](#)
226. From Wikipedia, the free encyclopedia. Jump to: navigation, search. This article is about the year 226. For the number 226, see 226 (number). ...
en.wikipedia.org/wiki/226 - Cached - Similar

[COS 226, Fall 2010: Home Page](#)
Princeton COS 226: Data Structures and Algorithms. ... Computer Science 226. Algorithms and Data Structures Fall 2010 ...
www.princeton.edu/~cos226/ - Cached - Similar

[Images for 226 - Report images](#)

RAP226



Done

<http://www.princeton.edu/~cos226>

Resources (web)

The screenshot shows a Google search results page with the following details:

- Search Bar:** 226 - Google Search
- Google Logo:** Google
- Search Query:** 226
- Search Tools:** Web (selected), Maps, Images, Videos, Books, More, Search tools
- Result Count:** About 175,000,000 results (0.15 seconds)
- First Result:**
 - Title:** COS 226, Spring 2014: Home Page
 - URL:** www.princeton.edu/~cos226/
 - Description:** Princeton University
If you have not taken COS 126 or COS 217 but want to place into COS 226, email Josh Hug. If you are unable to enroll because the only precepts you can attend ...
 - Links:** Assignments - Lectures - Exercises
- Second Result:**
 - Title:** 226 - Wikipedia, the free encyclopedia
 - URL:** en.wikipedia.org/wiki/226
 - Description:** Year 226 (CCXXVI) was a common year starting on Sunday (link will display the full calendar) of the Julian calendar. At the time, it was known as the Year of the ...
- Third Result:**
 - Title:** Area codes 519 and 226 - Wikipedia, the free encyclopedia
 - URL:** en.wikipedia.org/wiki/Area_codes_519_and_226
 - Description:** It is mostly bounded by area code 905, except for Simcoe County which is bordered by 705. It was overlaid with the new area code 226 on October 21, 2006, ...
 - Links:** History - Dialing areas - See also - References
- Fourth Result:**
 - Title:** Route 226 - King County Metro Transit
 - URL:** metro.kingcounty.gov/schedules/226/
 - Description:** Mar 28, 2013 - You are in: Travel Options >; Bus >; Schedules >; Route 226. Fares & ORCA · Trip Planner · Go to Metro Online's Home Page ...
 - Links:** To Overlake, Eastgate - Route Map - Saturday - Sunday

<http://www.princeton.edu/~cos226>

Resources (web)

The screenshot shows a Google search results page for the query "algorithms". The search bar at the top contains "algorithms". Below the search bar, the "Web" tab is selected, along with other options like Images, Maps, Shopping, News, More, and Search tools. A gear icon for settings is also present. The search results section displays approximately 22,700,000 results found in 0.13 seconds. The first result is a link to the Wikipedia page on algorithms, followed by a snippet of text explaining what an algorithm is. Subsequent results include links to the book "Algorithms, 4th Edition" by Robert Sedgewick and Kevin Wayne, a HowStuffWorks article on computer algorithms, and a Coursera course on algorithms.

algorithms - Google Search

Google algorithms

Web Images Maps Shopping News More Search tools

About 22,700,000 results (0.13 seconds)

[Algorithm - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/Algorithm

In mathematics and computer science, an **algorithm** is a step-by-step procedure for calculations. **Algorithms** are used for calculation, data processing, and ...

[List of algorithms - Algorithm examples - Automated reasoning - Euclidean](#)

[List of algorithms - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/List_of_algorithms

The following is a list of algorithms along with one-line descriptions for each.

[Algorithms, 4th Edition by Robert Sedgewick and Kevin Wayne](#)
algs4.cs.princeton.edu/

The textbook **Algorithms**, 4th Edition by Robert Sedgewick and Kevin Wayne surveys the most important **algorithms** and data structures in use today. The broad ...

[HowStuffWorks "What is a "computer algorithm"?"](#)
computer.howstuffworks.com/question717.htm

That's where computer **algorithms** come in. The **algorithm** is the basic technique used to get the job done. Let's follow an example to help get an understanding ...

[Algorithms, Part I | Coursera](#)
https://www.coursera.org/course/algs4partI

Algorithms, Part I is a free online class taught by Kevin Wayne and Robert Sedgewick of Princeton University.

Where to get help?

Piazza discussion forum.

- Low latency, low bandwidth.
- Mark solution-revealing questions as private.

piazza

<http://piazza.com/princeton/spring2014/cos226>

Office hours.

- High bandwidth, high latency.
- See web for schedule.



<http://www.princeton.edu/~cos226>


Computing laboratory.

- Undergrad lab TAs in Friend 017.
- For help with debugging.
- See web for schedule.




<http://www.princeton.edu/~cos226>

Where not to get help?



<http://world.edu/academic-plagiarism>



<http://www.youtube.com/watch?v=FT4NOe4vtoM>

What's ahead?

Lecture 1. [today] Union find.

Lecture 2. [Wednesday] Analysis of algorithms.

Flipped lecture 1. [Wednesday] Watch video beforehand.

Precept 1. [Thursday/Friday] Meets this week.



Exercise 1. Due via Bb submission at 11pm on Sunday.

Assignment 1. Due via electronic submission at 11pm on Tuesday.

protip: start early

Right course? See me.

Placed out of COS 126? Review Sections 1.1–1.2 of Algorithms 4/e.

Not registered? Go to any precept this week.


Change precept? Use SCORE.



see Colleen Kenny-McGinley
in CS 210 if the only precepts
you can attend are closed

Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE



ROBERT SEDGEWICK | KEVIN WAYNE

<http://algs4.cs.princeton.edu>

1.5 UNION-FIND

- ▶ *dynamic connectivity*
- ▶ *quick find*
- ▶ *quick union*
- ▶ *improvements*
- ▶ *applications*

Subtext of today's lecture (and this course)

Steps to developing a usable algorithm.

- Model the problem.
- Find an algorithm to solve it.
- Fast enough? Fits in memory?
- If not, figure out why not.
- Find a way to address the problem.
- Iterate until satisfied.

The scientific method.

Mathematical analysis.

Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<http://algs4.cs.princeton.edu>

1.5 UNION-FIND

- ▶ *dynamic connectivity*
- ▶ *quick find*
- ▶ *quick union*
- ▶ *improvements*
- ▶ *applications*

Dynamic connectivity problem

Given a set of N objects, support two operation:

- Connect two objects.
- Is there a path connecting the two objects?

connect 4 and 3

connect 3 and 8

connect 6 and 5

connect 9 and 4

connect 2 and 1

are 0 and 7 connected? ✗

are 8 and 9 connected? ✓


connect 5 and 0

connect 7 and 2

connect 6 and 1

connect 1 and 0

are 0 and 7 connected? ✓



A larger connectivity example

Q. Is there a path connecting p and q ?



A. Yes.

Modeling the objects

Applications involve manipulating objects of all types.

- Pixels in a digital photo.
- Computers in a network.
- Friends in a social network.
- Transistors in a computer chip.
- Elements in a mathematical set.
- Variable names in a Fortran program.
- Metallic sites in a composite system.

When programming, convenient to name objects 0 to $N - 1$.

- Use integers as array index.
- Suppress details not relevant to union-find.




can use symbol table to translate from site names to integers: stay tuned (Chapter 3)

Modeling the connections

We assume "is connected to" is an equivalence relation:

- Reflexive: p is connected to p .
- Symmetric: if p is connected to q , then q is connected to p .
- Transitive: if p is connected to q and q is connected to r ,
then p is connected to r .

Connected component. Maximal set of objects that are mutually connected.



{ 0 } { 1 4 5 } { 2 3 6 7 }


3 connected components

Implementing the operations


Find. In which component is object p ?

Connected. Are objects p and q in the same component?

Union. Replace components containing objects p and q with their union.



union(2, 5)
→



{ 0 } { 1 4 5 } { 2 3 6 7 }
3 connected components

{ 0 } { 1 2 3 4 5 6 7 }
2 connected components

Union-find data type (API)

Goal. Design efficient data structure for union-find.

- Number of objects N can be huge.
- Number of operations M can be huge.
- Union and find operations may be intermixed.

```
public class UF
```

```
UF(int N)
```

*initialize union-find data structure
with N singleton objects (0 to $N - 1$)*

```
void union(int p, int q)
```

add connection between p and q

```
int find(int p)
```

component identifier for p (0 to $N - 1$)

```
boolean connected(int p, int q)
```

are p and q in the same component?

```
public boolean connected(int p, int q)  
{ return find(p) == find(q); }
```

1-line implementation of connected()

Dynamic-connectivity client

- Read in number of objects N from standard input.
- Repeat:
 - read in pair of integers from standard input
 - if they are not yet connected, connect them and print out pair

```
public static void main(String[] args)
{
    int N = StdIn.readInt();
    UF uf = new UF(N);
    while (!StdIn.isEmpty())
    {
        int p = StdIn.readInt();
        int q = StdIn.readInt();
        if (!uf.connected(p, q))
        {
            uf.union(p, q);
            StdOut.println(p + " " + q);
        }
    }
}
```

% more tinyUF.txt

10	
4	3
3	8
6	5
9	4
2	1
8	9
5	0
7	2
6	1
1	0
6	7

already connected

```
graph TD; 0 --> 1; 0 --> 2; 0 --> 5; 0 --> 6; 0 --> 7; 1 --> 2; 3 --> 8; 4 --> 9; 5 --> 6; 6 --> 1; 6 --> 7; 7 --> 6; 8 --> 9; 9 --> 4;
```

Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<http://algs4.cs.princeton.edu>

1.5 UNION-FIND

- ▶ *dynamic connectivity*
- ▶ *quick find*
- ▶ *quick union*
- ▶ *improvements*
- ▶ *applications*

Quick-find [eager approach]

Data structure.


- Integer array `id[]` of length N .
- Interpretation: `id[p]` is the id of the component containing p .

if and only if

↓

	0	1	2	3	4	5	6	7	8	9
<code>id[]</code>	0	1	1	8	8	0	0	1	8	8

0, 5 and 6 are connected
1, 2, and 7 are connected
3, 4, 8, and 9 are connected



Quick-find [eager approach]

Data structure.

- Integer array $\text{id}[]$ of length N .
- Interpretation: $\text{id}[p]$ is the id of the component containing p .

	0	1	2	3	4	5	6	7	8	9
$\text{id}[]$	0	1	1	8	8	0	0	1	8	8

Find. What is the id of p ?

$\text{id}[6] = 0; \text{id}[1] = 1$
6 and 1 are not connected

Connected. Do p and q have the same id?

Union. To merge components containing p and q , change all entries whose id equals $\text{id}[p]$ to $\text{id}[q]$.

	0	1	2	3	4	5	6	7	8	9
$\text{id}[]$	1	1	1	8	8	1	1	1	8	8

problem: many values can change

after union of 6 and 1

Quick-find demo



0

1

2

3

4

5

6


7

8

9

	0	1	2	3	4	5	6	7	8	9
id[]	0	1	2	3	4	5	6	7	8	9

Quick-find demo



	0	1	2	3	4	5	6	7	8	9
id[]	1	1	1	8	8	1	1	1	8	8

Quick-find: Java implementation

```
public class QuickFindUF
{
    private int[] id;

    public QuickFindUF(int N)
    {
        id = new int[N];
        for (int i = 0; i < N; i++)
            id[i] = i;
    }

    public int find(int p)
    {
        return id[p];
    }

    public void union(int p, int q)
    {
        int pid = id[p];
        int qid = id[q];
        for (int i = 0; i < id.length; i++)
            if (id[i] == pid) id[i] = qid;
    }
}
```

set id of each object to itself
(N array accesses)

return the id of p
(1 array access)

change all entries with $\text{id}[p]$ to $\text{id}[q]$
(at most $2N + 2$ array accesses)

Quick-find is too slow

Cost model. Number of array accesses (for read or write).

algorithm	initialize	union	find	connected
quick-find	N	N	1	1

order of growth of number of array accesses

Union is too expensive. It takes N^2 array accesses to process a sequence of N union operations on N objects.

quadratic




Quadratic algorithms do not scale

Rough standard (for now).

- 10^9 operations per second.
- 10^9 words of main memory.
- Touch all words in approximately 1 second.

a truism (roughly
since 1950!)




Ex. Huge problem for quick-find.

- 10^9 union commands on 10^9 objects.
- Quick-find takes more than 10^{18} operations.
- 30+ years of computer time!

Quadratic algorithms don't scale with technology.

- New computer may be 10x as fast.
- But, has 10x as much memory ⇒ want to solve a problem that is 10x as big.
- With quadratic algorithm, takes 10x as long!



Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<http://algs4.cs.princeton.edu>

1.5 UNION-FIND


- ▶ *dynamic connectivity*
- ▶ *quick find*
- ▶ *quick union*
- ▶ *improvements*
- ▶ *applications*

Quick-union [lazy approach]

Data structure.

- Integer array `id[]` of length N .
- Interpretation: `id[i]` is parent of i . keep going until it doesn't change
(algorithm ensures no cycles)
- Root of i is `id[id[id[...id[i]...]]]`.

	0	1	2	3	4	5	6	7	8	9
<code>id[]</code>	0	1	9	4	9	6	6	7	8	9



parent of 3 is 4


root of 3 is 9

Quick-union [lazy approach]

Data structure.

- Integer array $\text{id}[]$ of length N .
- Interpretation: $\text{id}[i]$ is parent of i .
- Root of i is $\text{id}[\text{id}[\text{id}[\dots \text{id}[i]\dots]]]$.

0	1	2	3	4	5	6	7	8	9	
$\text{id}[]$	0	1	9	4	9	6	6	7	8	9



Find. What is the root of p ?

Connected. Do p and q have the same root?

root of 3 is 9


root of 5 is 6

3 and 5 are not connected

Union. To merge components containing p and q , set the id of p 's root to the id of q 's root.

0	1	2	3	4	5	6	7	8	9	
$\text{id}[]$	0	1	9	4	9	6	6	7	8	6

↑
only one value changes




Quick-union demo



	0	1	2	3	4	5	6	7	8	9
id[]	0	1	2	3	4	5	6	7	8	9

Quick-union demo



Quick-union: Java implementation

```
public class QuickUnionUF
{
```

```
    private int[] id;
```

```
    public QuickUnionUF(int N)
```

```
    {
```

```
        id = new int[N];
```

```
        for (int i = 0; i < N; i++) id[i] = i;
```

```
}
```

```
    public int find(int i)
```

```
    {
```

```
        while (i != id[i]) i = id[i];
```

```
        return i;
```

```
}
```

set id of each object to itself
(N array accesses)

```
    public void union(int p, int q)
```

```
    {
```

```
        int i = find(p);
```

```
        int j = find(q);
```

```
        id[i] = j;
```

```
}
```

chase parent pointers until reach root
(depth of i array accesses)

change root of p to point to root of q
(depth of p and q array accesses)

Quick-union is also too slow

Cost model. Number of array accesses (for read or write).

algorithm	initialize	union	find	connected
quick-find	N	N	1	1
quick-union	N	N †	N	N

← worst case

† includes cost of finding roots

Quick-find defect.

- Union too expensive (N array accesses).
- Trees are flat, but too expensive to keep them flat.

Quick-union defect.

- Trees can get tall.
- Find/connected too expensive (could be N array accesses).

Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<http://algs4.cs.princeton.edu>


1.5 UNION-FIND

- ▶ *dynamic connectivity*
- ▶ *quick find*
- ▶ *quick union*
- ▶ ***improvements***
- ▶ *applications*

Improvement 1: weighting

Weighted quick-union.

- Modify quick-union to avoid tall trees.
- Keep track of size of each tree (number of objects).
- Balance by linking root of smaller tree to root of larger tree.




Weighted quick-union demo




	0	1	2	3	4	5	6	7	8	9
id[]	0	1	2	3	4	5	6	7	8	9

Weighted quick-union demo




Quick-union and weighted quick-union example

quick-union



average distance to root: 5.11

weighted



average distance to root: 1.52

Quick-union and weighted quick-union (100 sites, 88 union() operations)

Weighted quick-union: Java implementation

Data structure. Same as quick-union, but maintain extra array $sz[i]$ to count number of objects in the tree rooted at i .

Find/connected. Identical to quick-union.

Union. Modify quick-union to:

- Link root of smaller tree to root of larger tree.
- Update the $sz[]$ array.

```
int i = find(p);
int j = find(q);
if (i == j) return;
if (sz[i] < sz[j]) { id[i] = j; sz[j] += sz[i]; }
else                  { id[j] = i; sz[i] += sz[j]; }
```


Weighted quick-union analysis

Running time.

- Find: takes time proportional to depth of p .
- Union: takes constant time, given roots.

Proposition. Depth of any node x is at most $\lg N$.

\lg = base-2 logarithm



Weighted quick-union analysis

Running time.

- Find: takes time proportional to depth of p .
- Union: takes constant time, given roots.


\lg = base-2 logarithm

Proposition. Depth of any node x is at most $\lg N$.

Pf. What causes the depth of object x to increase?

Increases by 1 when tree T_1 containing x is merged into another tree T_2 .

- The size of the tree containing x at least doubles since $|T_2| \geq |T_1|$.
- Size of tree containing x can double at most $\lg N$ times. Why?



Weighted quick-union analysis

Running time.

- Find: takes time proportional to depth of p .
- Union: takes constant time, given roots.

Proposition. Depth of any node x is at most $\lg N$.


algorithm	initialize	union	find	connected
quick-find	N	N	1	1
quick-union	N	N^{\dagger}	N	N
weighted QU	N	$\lg N^{\dagger}$	$\lg N$	$\lg N$

\dagger includes cost of finding roots

- Q.** Stop at guaranteed acceptable performance?
A. No, easy to improve further.


Improvement 2: path compression

Quick union with path compression. Just after computing the root of p , set the `id[]` of each examined node to point to that root.




Improvement 2: path compression

Quick union with path compression. Just after computing the root of p , set the `id[]` of each examined node to point to that root.




Improvement 2: path compression

Quick union with path compression. Just after computing the root of p , set the `id[]` of each examined node to point to that root.




Improvement 2: path compression

Quick union with path compression. Just after computing the root of p , set the `id[]` of each examined node to point to that root.



Improvement 2: path compression

Quick union with path compression. Just after computing the root of p , set the `id[]` of each examined node to point to that root.



Bottom line. Now, `find()` has the side effect of compressing the tree.

Path compression: Java implementation

Two-pass implementation: add second loop to find() to set the id[] of each examined node to the root.

Simpler one-pass variant (path halving): Make every other node in path point to its grandparent.

```
public int find(int i)
{
    while (i != id[i])
    {
        id[i] = id[id[i]]; ← only one extra line of code !
        i = id[i];
    }
    return i;
}
```

In practice. No reason not to! Keeps tree almost completely flat.

Weighted quick-union with path compression: amortized analysis

Proposition. [Hopcroft-Ulman, Tarjan] Starting from an empty data structure, any sequence of M union–find ops on N objects makes $\leq c(N + M \lg^* N)$ array accesses.

- Analysis can be improved to $N + M \alpha(M, N)$.
- Simple algorithm with fascinating mathematics.

N	$\lg^* N$
1	0
2	1
4	2
16	3
65536	4
2^{65536}	5

iterated \lg function

Linear-time algorithm for M union–find ops on N objects?

- Cost within constant factor of reading in the data.
- In theory, WQUPC is not quite linear.
- In practice, WQUPC is linear.

Amazing fact. [Fredman-Saks] No linear-time algorithm exists.

in "cell-probe" model of computation

Summary

Key point. Weighted quick union (and/or path compression) makes it possible to solve problems that could not otherwise be addressed.

algorithm	worst-case time
quick-find	$M N$
quick-union	$M N$
weighted QU	$N + M \log N$
QU + path compression	$N + M \log N$
weighted QU + path compression	$N + M \lg^* N$

order of growth for M union-find operations on a set of N objects

Ex. [10⁹ unions and finds with 10⁹ objects]

- WQUPC reduces time from 30 years to 6 seconds.
- Supercomputer won't help much; good algorithm enables solution.

Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE




<http://algs4.cs.princeton.edu>

1.5 UNION-FIND

- ▶ *dynamic connectivity*
- ▶ *quick find*
- ▶ *quick union*
- ▶ *improvements*
- ▶ *applications*

Union-find applications


- Percolation.
- Games (Go, Hex).
- ✓ Dynamic connectivity.
 - Least common ancestor.
 - Equivalence of finite state automata.
 - Hoshen-Kopelman algorithm in physics.
 - Hinley-Milner polymorphic type inference.
 - Kruskal's minimum spanning tree algorithm.
 - Compiling equivalence statements in Fortran.
 - Morphological attribute openings and closings.
 - Matlab's bwlabel() function in image processing.



Percolation

An abstract model for many physical systems:

- N -by- N grid of sites.
- Each site is open with probability p (and blocked with probability $1 - p$).
- System **percolates** iff top and bottom are connected by open sites.



Percolation


An abstract model for many physical systems:

- N -by- N grid of sites.
- Each site is open with probability p (and blocked with probability $1 - p$).
- System **percolates** iff top and bottom are connected by open sites.


model	system	vacant site	occupied site	percolates
electricity	material	conductor	insulated	conducts
fluid flow	material	empty	blocked	porous
social interaction	population	person	empty	communicates

Likelihood of percolation


Depends on grid size N and site vacancy probability p .






p low (0.4)
does not percolate



p medium (0.6)
percolates?



p high (0.8)
percolates




Percolation phase transition

When N is large, theory guarantees a sharp threshold p^* .


- $p > p^*$: almost certainly percolates.
- $p < p^*$: almost certainly does not percolate.

Q. What is the value of p^* ?



Monte Carlo simulation

- Initialize all sites in an N -by- N grid to be blocked.
- Declare random sites open until top connected to bottom.
- Vacancy percentage estimates p^* .




Dynamic connectivity solution to estimate percolation threshold

Q. How to check whether an N -by- N system percolates?

A. Model as a **dynamic connectivity** problem and use **union-find**.

$N = 5$



open site




blocked site

Dynamic connectivity solution to estimate percolation threshold

Q. How to check whether an N -by- N system percolates?

- Create an object for each site and name them 0 to $N^2 - 1$.


$N = 5$



open site





blocked site



Dynamic connectivity solution to estimate percolation threshold

Q. How to check whether an N -by- N system percolates?

- Create an object for each site and name them 0 to $N^2 - 1$.
- Sites are in same component iff connected by open sites.



open site



blocked site

Dynamic connectivity solution to estimate percolation threshold

Q. How to check whether an N -by- N system percolates?

- Create an object for each site and name them 0 to $N^2 - 1$.
- Sites are in same component iff connected by open sites.
- Percolates iff any site on bottom row is connected to any site on top row.

brute-force algorithm: N^2 calls to `connected()`



open site



blocked site

Dynamic connectivity solution to estimate percolation threshold

Clever trick. Introduce 2 virtual sites (and connections to top and bottom).

- Percolates iff virtual top site is connected to virtual bottom site.

more efficient algorithm: only 1 call to connected()




open site

blocked site

Dynamic connectivity solution to estimate percolation threshold

Q. How to model opening a new site?




Dynamic connectivity solution to estimate percolation threshold

Q. How to model opening a new site?

A. Mark new site as open; connect it to all of its adjacent open sites.


up to 4 calls to `union()`



open site



blocked site




Percolation threshold

Q. What is percolation threshold p^* ?

A. About 0.592746 for large square lattices.

constant known only via simulation



Fast algorithm enables accurate answer to scientific question.

Subtext of today's lecture (and this course)

Steps to developing a usable algorithm.

- Model the problem.
- Find an algorithm to solve it.
- Fast enough? Fits in memory?
- If not, figure out why.
- Find a way to address the problem.
- Iterate until satisfied.

The scientific method.

Mathematical analysis.